

Parallel Algorithms

Umut Acar
Carnegie Mellon University
umutaco@gmail.com
umut@cmu.edu

July 9, 2018

1 Language Based Cost Models

- Abstract
- Truthful

Key Assumption Throughout: Pure Functional Programs

Shared state - Two programs that share the same data and each run ten lines of code (each) that rely on that data, then to prove correctness need to check 2^{20} combinations or 10^{100} computations. Pure functional programming does not have shared state, and therefore avoids this problem.

1.1 λ -Calculus

Cost Semantics is simple with a little twist.

1. Inherently parallel ($e_1 \parallel e_2$)
2. Twist for cost - Work (additive) and Span (max)

Example:

```
fun f x =  
  if x ≤ 1 then  
    x  
  else  
    let (a,b) = (f(x-1)∥f(x-2))  
    in  
      a+b  
    end
```

Work:

$$\begin{aligned} W(n) &= 1 \text{ if } n \leq 1 \\ &= W(n-1) + W(n-2) + 1 \end{aligned}$$

Span:

$$\begin{aligned} S(n) &= 1 \text{ if } n \leq 1 \\ &= \text{Max}(S(n-1), S(n-2)) + 1 \end{aligned}$$

Good parallel algorithms have work as close to the sequential version as possible and have low span ($O(\log^d(n))$).

1.2 Provably Efficient Implementations - Bounded Implementations

Someone provides an algorithm: $\text{Alg}(W, S)$, and our goal is to provide a programming language that proves the efficiency.

What we need:

- The result should be the same as sequential.
- $T_p \geq \frac{W}{p}$ and $T_p \geq S$ are lower bounds.
- Upper bound?

Note: T_p^{OPT} is NP-Complete

Schedule:

	<u>1</u>	<u>2</u>
1	a	
2	b	c
3	d	e
4	f	
5	g	

Brent's Theorem - 1974

W_1 is the number of work at level 1

$$W = \sum_{i=1}^S W_i$$

$$T_p = \sum_{i=1}^S \left\lceil \frac{W_i}{p} \right\rceil = \sum_{i=1}^S \left\lfloor \frac{W_i}{p} \right\rfloor + 1 = \left(\sum_{i=1}^S \left\lfloor \frac{W_i}{p} \right\rfloor \right) + S \leq \left(\sum_{i=1}^S \frac{W_i}{p} \right) + S = \frac{W}{p} + S$$

Implications:

$$T_p \leq \frac{W}{p} + S \leq 2 * OPT$$

$$\left(T_p \geq \frac{W}{p}, T_p \geq S \right) \implies T_p \geq \max\left(\frac{W}{p}, S\right) \implies OPT \geq \max\left(\frac{W}{p}, S\right)$$

Greedy:

If \exists idle processor, and \exists unprocessed vertex then assign

$$T_p \leq \frac{W}{p} + S \frac{(p-1)}{p}$$

Think of it as two buckets, a work bucket and an idle bucket. Every time a computation is made then place a coin into the work bucket. If there is an idle processor, then put a coin in the idle bucket. The total coins in the work bucket is W and the total number of coins placed into the idle bucket at each step is $(p-1)$ and there are a total of S steps, so the total coins that can be in the idle bucket is $(p-1)S$ and $T_p = \frac{\text{total coins}}{p} \leq \frac{W+(p-1)S}{p}$

1.3 Load Balancing (Work Stealing)

But how can a work queue give out more than one task at a time? It can't, so need a queue for each processor (can think of this as a call stack).

But distributing from one processor queue to the other is a concurrent algorithm, not a parallel algorithm, and is extremely difficult to keep straight.

In a cost model for λ -calculus the run-time system does this for you.

$$E[T_p] \leq \frac{W}{p} + S \text{ In expectation} \leftarrow \text{TRUTHFUL!}$$